

Value Pluralism, Lyotard's Metanarrative Problem and the Foundations of AI Alignment

The direction of AI alignment appears to be straightforward. *What exactly is AI aiming to align itself to?* is a seemingly simple question that permeates most AI alignment difficulties. The majority of alignment literature treats this as a technical problem, not a philosophical one. The assumption being that if we can correctly specify our values, teach them or have AI optimise towards them, the problems faced by AI safety specialists might dissipate. However, this, in itself, carries hidden assumptions about the nature of values, knowledge, and human universality.

AI alignment is, at least in the mainstream literature, understood as the project of ensuring that AI systems pursue goals and exhibit behaviours that are beneficial to humans. It breaks down into four main paradigms: value learning, RHLF, constitutional AI, and corrigibility (BlueDot Impact, 2025). In value learning, the AI system infers human values from behaviour. In RHLF, human raters train the model toward preferred outputs (Soares, 2016). Constitutional AI involves encoding explicit normative principles into the model and corrigibility requires ensuring the AI defers to human correction (Bai et al., 2022). However, across all these paradigms, the question of whose values (which humans' values) they espouse and what counts as beneficial is either assumed or treated as a downstream problem, rather than a foundational one.

AI alignment discourse, despite its technical sophistication, rests on an unreflective modernist metanarrative. The assumption is that there exists a coherent, discoverable, and sufficiently universal set of "human values" toward which AI systems can be directed. Drawing on Jean-François Lyotard's *The Postmodern Condition*, this paper will show how alignment thinking reproduces the structure of a metanarrative, apply Lyotard's critique to expose the political and epistemological stakes of that move and explore what a post-metanarrative approach to alignment might demand (Lyotard, 1979).

Lyotard and The Postmodern Condition

In *The Postmodern Condition*, Lyotard argues that we are entering a condition in which the grand, legitimating stories that once justified knowledge and social organisation are losing their credibility (Lyotard, 1979). This is situated against the backdrop of post-1968 disillusionment, the collapse of revolutionary optimism, the fragmentation of the political left, and the exposure of totalitarian uses of "universal" ideologies (Gratton, 2018). Lyotard defines *un grand récit* as a story that legitimates knowledge, science, or political practice by appealing to a grand subject (humanity, reason, the proletariat, the spirit) and a grand telos (liberation, progress, emancipation) (Lyotard, 1979). He identifies two key examples: the Enlightenment

metanarrative where knowledge produces freedom and science liberates humanity from ignorance and superstition and the Hegelian metanarrative where knowledge is the self-realisation of Spirit and the university unifies all knowledge into a coherent whole (Gratton, 2018). Both these metanarratives share the same structure: a universal subject (humanity, reason), a progressive directionality (toward the good, the true, the free) and a claim to transcend local, particular, or interested perspectives (Peters, 1997). Lyotard defines the “*postmodern as incredulity toward metanarratives*” (Lyotard, 1979). This is a historical and epistemological claim. These stories have lost their legitimating force because internal contradictions have been exposed (e.g. progress produced ecological catastrophe); and no single narrative can account for the multiplicity and incommensurability of language games (a term borrowed from Wittgenstein) (Lyotard, 1979).

Lyotard draws on Wittgenstein's concept of language games (Biletzki, 2021), forms of life with their own internal rules, criteria of validity, and modes of meaning-making (Lyotard, 1979). He claims that different language games (scientific, ethical, aesthetic, political) are incommensurable (Lyotard, 1979). There is no master language or meta-criterion that can adjudicate between them (Gratton, 2018). He defines the *différend* as the irreducible conflict between parties whose idioms are so different that a wrong suffered by one cannot be expressed in the terms available to the other (Lyotard, 1979). This is a structural injustice built into the very architecture of discourse (Lyotard, 1979). Any attempt to impose a single framework of value across incommensurable communities is itself an act of violence (a silencing of the *différend*) (Lyotard, 1979). We can also see links to Foucault in his analysis of knowledge, power and legitimacy (Gutting & Oksala, 2026). Lyotard notes that metanarratives always serve particular interests while claiming universality (Gratton, 2018). He points out the performativity criterion whereby in late capitalist societies, knowledge is increasingly legitimated not by its truth but by its efficiency and utility (Lyotard, 1979). This produces a new threat. The technocratic metanarrative of optimisation, the idea that the best system is the most efficient one (Lyotard, 1979). This critique will become central to the AI alignment critique posed in this paper.

AI Alignment as a Metanarrative

Applying Lyotard's analysis to AI alignment discourse reveals that there is an implicit *grand récit* at play. This story has the following structure:

- There exists a set of human values
- These values can be identified, aggregated, and represented
- AI systems can be directed toward these values
- This is what it means for AI to be beneficial to humanity

This story has the exact structure of a Lyotardian metanarrative. The universal subject is humanity or rather human preferences. The grand telos is beneficial outcomes, human flourishing, and safety. The legitimating function is the justification of the vast technical and economic apparatus of AI development.

However, the way alignment frameworks construct their universal subject is problematic. In RHLF, human feedback is gathered from a narrow pool of raters (predominantly English-speaking, Western, relatively affluent) and these preferences are then generalised as human values (Casper et al., 2023). With Constitutional AI, Anthropic's approach encodes explicit principles, but we must ask who authors those principles and by what authority do they do so (Bai et al., 2022). Also, aggregative approaches assume that diverse preferences can be summed or averaged into a coherent collective preference (Bao et al., 2026). Arrow's Impossibility Theorem already shows the mathematical limits of this within a single community (Morreau, 2025). For Lyotard, the communities themselves are incommensurable.

We may also consider the performativity trap. The danger here is that value is reduced to efficiency. For most alignment frameworks, beneficial is defined in terms of measurable outcomes: reduced harm scores, higher human preference ratings, lower perplexity on safety benchmarks (Seneque et al., 2024; Jalan et al., 2026). This is exactly the technocratic metanarrative Lyotard described: the replacement of ethical legitimation with technical optimisation. We might think of Goodhart's law problem as a postmodern symptom: when a measure becomes a target, it ceases to be a good measure, so the metric colonises the value (Mattson et al., 2021). Discussions of an alignment tax illustrate this. Commercial AI labs face pressure to define alignment in ways that are compatible with deployment and profit so "safe" becomes "safe enough to ship" (Young, 2026).

Alignment is the telos of this progress narrative. Most alignment discourse implicitly suggests that we are moving toward increasingly safe, increasingly aligned systems. Structurally, this is a metanarrative of emancipation. Humanity will eventually be freed from the risk of misaligned AI through the advance of alignment science. This narrative forecloses the possibility that alignment itself is the wrong frame. Safety is not a destination but a permanent negotiation. There is an eschatological structure inherent in some alignment thinking (e.g., the "superintelligence" discourse of Nick Bostrom). It describes singular, catastrophic event horizon that gives urgency and shape to the entire project (Bostrom, 2014). Again, this is grand narrative par excellence.

The Différend in AI Alignment

We can apply Lyotard's concept of the *différend* to alignment's exclusions (Lyotard, 1979). This requires analysing which communities cannot express their grievances in the idiom of human values and preference optimisation. For instance, communities whose primary harms are relational, spiritual, or cosmological rather than material, future generations, who cannot participate in current preference aggregation and

non-human beings (animals, ecosystems), whose exclusion from human values is built into the frame are all constituents of the *différend*. Their voices are not often heard by alignment discourse and engineering.

Again, safety is itself a language game with its own rules, criteria, and exclusions. What counts as safe in AI alignment discourse (factual accuracy, non-toxicity, refusal of harmful requests, CSAM filtering, etc.) and what does not register as unsafe within this language game (economic displacement, surveillance infrastructure, cultural homogenisation, epistemic dependence etc.) is very revealing (Dahlgren Lindström et al., 2025). It pays homage to a very particular set of values and communities harmed in ways that do not fit the alignment language game cannot press their case. While we may think of relative misalignment (when an AI system is aligned in one community but fails for another) (LaCroix, 2026), as a failure of representation, this analysis reveals that it may instead be a structural *différend*.

Rescuing AI Alignment from the Postmodern Critique

The postmodern critique does not dissolve the alignment problem. AI systems must be built and deployed now, under conditions of irreducible value pluralism. Proceduralism, where communities agree on fair processes rather than outcomes, is an inadequate response because, according to Lyotard, procedures are not neutral; they encode their own language games and exclusions. The goal is justice, which requires attending to the *différend*, we cannot resolve it.

Instead, we must reframe alignment as permanent, agonistic negotiation. It is not as simple as a technical problem to be solved. To do this, we need to build contestability into systems as a first-order value. It cannot be a just safety certification; it must be an ongoing process. We must also treat every deployment as provisional and revisable, not a finished product and institutionalise adversarial review of alignment criteria, rather than just reviewing model outputs. We must reconceptualise AI alignment: alignment without secure foundations, without a final destination, and without a universal subject. Of course, the prevailing difficulty with this is that this approach may be structurally incompatible with the commercial logic driving AI development (which is itself a political conclusion, not a failure of imagination).

Conclusion

Alignment discourse depends on a metanarrative of universal human values that Lyotard's postmodern critique exposes as philosophically untenable and politically dangerous. The *différend* shows us not just who is excluded but why they cannot speak. It is structural, not accidental. I suggest reframing AI alignment as a permanent negotiation. This goes beyond an academic exercise. The alignment

problem is perhaps one of the most consequential design problems in human history. Getting the philosophical foundations wrong has practical consequences. We might easily end up with systems that are aligned in name while encoding the values of a narrow, powerful elite. Postmodern thought offers indispensable tools for interrogating these foundations. In response to the opening question (*what exactly is AI aiming to become aligned to?*), the answer this paper proposes is that is we do not yet know, we cannot fully know, and the pretence that we do is itself a form of power. The task at hand is not to answer the question once and for all, but to keep it open and to build AI systems that remain answerable to the question rather than systems that presume to have resolved it.

Bibliography

- Bai, Y., Kadavath, S., Kundu, S., Askill, A., Kernion, J., Jones, A., Chen, A., Goldie, A., Mirhoseini, A., Mckinnon, C., Chen, C., Olsson, C., Olah, C., Hernandez, D., Drain, D., Ganguli, D., Li, D., Tran-Johnson, E., Perez, E., ... Kaplan, J. (2022). *Constitutional AI: Harmlessness from AI Feedback*.
- Bao, H., Huang, Y., Wang, X., Zhang, Z., Zhou, Y., Yang, C., Zhang, X., & Ye, Y. (2026). *AI Alignment Breaks at the Edge*. <https://arxiv.org/pdf/2602.20042v2>
- Biletzki, A. (2021). *Ludwig Wittgenstein (Stanford Encyclopedia of Philosophy)*. <https://plato.stanford.edu/entries/wittgenstein/>
- BlueDot Impact. (2025). *Reinforcement Learning from Human Feedback (RLHF): A Simple Explainer*. <https://blog.bluedot.org/p/reinforcement-learning-from-human>
- Bostrom, N. (2014). *Superintelligence: Paths, Dangers, Strategies*.
- Casper, S., Davies, X., Shi, C., Gilbert, T. K., Scheurer, J., Rando, J., Freedman, R., Korbak, T., Lindner, D., Freire, P., Wang, T., Marks, S., Segerie, C. R., Carroll, M., Peng, A., Christoffersen, P., Damani, M., Slocum, S., Anwar, U., ... Hadfield-Menell, D. (2023). Open Problems and Fundamental Limitations of Reinforcement Learning from Human Feedback. *Transactions on Machine Learning Research*, 2023.

- Dahlgren Lindström, A., Methnani, L., Krause, L., Ericson, P., de Rituerto de Troya, Í. M., Coelho Mollo, D., & Dobbe, R. (2025). Helpful, harmless, honest? Sociotechnical limits of AI alignment and safety through Reinforcement Learning from Human Feedback. *Ethics and Information Technology*, 27(2), 28. <https://doi.org/10.1007/S10676-025-09837-2>
- Gratton, P. (2018). *Jean François Lyotard* (Stanford Encyclopedia of Philosophy). Stanford Encyclopedia of Philosophy. <https://plato.stanford.edu/entries/lyotard/>
- Gutting, G., & Oksala, J. (2026). *Michel Foucault* (Stanford Encyclopedia of Philosophy). <https://plato.stanford.edu/entries/foucault/>
- Jalan, P., Abishethvarman, V., Chandna, B., & Naseem, U. (2026). Survey on LLM Safety: Attacks, Defenses, Alignment, Metrics, and Guardrails. *Machine Learning 2026* 115:6, 115(6), 130-. <https://doi.org/10.1007/S10994-026-07060-8>
- LaCroix, T. (2026). Relative Principals, Pluralistic Alignment, and the Structural Value Alignment Problem. *The 2026 ACM Conference on Fairness, Accountability, and Transparency (FAccT '26)*, June 25â•fi28, 2026, Montreal, QC, Canada, 1, 18. <https://doi.org/10.1145/3805689.3812420>
- Lyotard, J.-F. (1979). *Selections from: The Postmodern Condition: A Report on Knowledge*.
- Mattson, C., Bushardt, R. L., & Artino, A. R. (2021). "When a Measure Becomes a Target, It Ceases to be a Good Measure." *Journal of Graduate Medical Education*, 13(1), 2. <https://doi.org/10.4300/JGME-D-20-01492.1>
- Morreau, M. (2025). *Arrow's Theorem* (Stanford Encyclopedia of Philosophy). <https://plato.stanford.edu/entries/arrows-theorem/>
- Peters, M. (1997). Lyotard, Education, and the Problem of Capitalism in the Postmodern Condition. In *Educational Research Association*.
- Seneque, G., Ho, L.-H., Kuperman, A., Saeedi, N. E., & Molendijk, J. (2024). *ABC Align: Large Language Model Alignment for Safety & Accuracy*. <https://arxiv.org/pdf/2408.00307v1>
- Soares, N. (2016). *The Value Learning Problem In: Ethics for Artificial*.
- Young, R. (2026). *What Is the Alignment Tax?* <https://arxiv.org/pdf/2603.00047v2>