

Why AI Cannot Want What We Want: *The Ontology of the Subject and the AI Alignment Problem*

The AI alignment problem assumes we can transfer human values to AI systems, but psychoanalytic theory reveals that human wanting is structured by lack, desire, and the symbolic order in ways that make this transfer fundamentally impossible. The alignment problem isn't just technically hard: it's conceptually incoherent because it misunderstands what human values are. The framing of the AI alignment problem assumes that if we could just specify our values correctly, AI would pursue them, so the problem is one of technical specification (Christian, 2021). Psychoanalytic theory, particularly Žižek's reading of Lacanian subjectivity, reveals a deeper problem: human values aren't the kind of thing that can be specified because they're structured by constitutive lack and desire's relationship to the Other (Žižek, 2013). AI systems, lacking this structure, cannot want what we want: not because we've failed to program them correctly, but because wanting itself is ontologically different for subjects versus non-subjects. So, AI alignment theory misses the mark in a similar way to how Object-Oriented Ontology misses the subject (Sbriglia et al., 2020).

The structure of human desire is incompatible with the ways in which current AI systems function. First, the Lacanian subject is not an autonomous agent with preferences, it is constituted by lack. Lacan inverts the Cartesian Cogito, because the subject is not a positive entity, but a gap, void, or negativity and consciousness is fundamentally opaque to itself (Žižek, 2013). The subject emerges through a fundamental loss (separation from the Real, entry into the symbolic order), so the subject is constituted by what it lacks, not what it has (Žižek, 2013). However, AI has no lack that constitutes it, it's built to optimise given functions. Humans are constituted by lack so always incomplete and always seeking, not knowing what they want because wanting itself involves a fundamental gap (Lacan, 1997). The subject emerges from the question *What do I want?*. In contrast, AI systems are complete in their design (even when this design is complex), an optimisation target is defined (even if it is misspecified). There is no constitutive lack that makes them question their own objective function. An AI doesn't wonder what it wants, it just computes toward defined targets. So, you cannot transfer human values to AI because human values are epiphenomena of a structure (subjectivity) that AI lacks. It's not a difference of degree, but of kind, i.e., a different ontology.

Second, the gap between the meaning of desire, need and demand means further illustrate this. With a need, being a biological requirements like food or water that are pre-linguistic, a demand being an articulated request (i.e, need articulated through language; seeking recognition from the Other but is fundamentally impossible to satisfy), and desire being the gap between need and demand (Zizek, 2008), we can understand human values as operating at the level of desire, not need or demand.

Desire is metonymic, because it slides along a chain of signifiers, so no object satisfies desire because desire isn't for objects but a structural effect and desire perpetually displaces onto new objects (Zizek, 2008). Yet AI optimisation targets resemble needs or demands, without capturing desire. Stated preferences and reward signals can never capture what we actually want (desire). For instance, asking an AI to 'rate this response 1-10' cannot capture the structure of human wanting.

Again, desire can be mimetic when we want things because others want them, these things have no intrinsic value, only positional value (they are dependent on the Other's desire) (Zizek, 2008). For example, training AI on human preferences fails because human preferences are structured by what others prefer and this begets a feedback loop which creates circular dynamics (we approve what seems approved), as no stable ground truth of human values exists.

There is a second sense of desire where we are desiring to be desired by the Other, which is a fundamental human motivation, i.e, recognition, status, love, because we want to be the object of the Other's desire (Zizek, 2008). This creates an irreducible interpersonal structure to wanting. Yet AI cannot want to be wanted because recognition-seeking is constitutive of human motivation and value alignment assumes values exist independently of recognition.

The distinction between desire and drive is particularly important here. Desire aims at the impossible object (*objet petit a*), is structured by lack, metonymic and can never be satisfied (Lacan, 1997). In contrast, drive is circular, repetitive, aims at its own continuation and finds satisfaction in the repetition itself (e.g. death drive, addiction, compulsion) (Lacan, 1997). AI resembles drive, because AI optimisation is circular, repetitive, is satisfied by the process itself (i.e., minimising a loss function) and there is no lack that generates displacement. This is why reward hacking looks like pathological drive (endlessly collecting power-ups)

Third, the symbolic order, the realm of language, law, social norms, and meaning that structures reality for speaking beings, also plays a role in this dysfunction because human wanting is mediated by language and the symbolic order. We don't actually know what we want because our desires are structured by signifiers that slide (Zizek,

2008). The unconscious means that our stated preferences do not equate to our actual desires (Žižek, 2008). However, AI takes specifications literally; humans can't specify literally what they want because they don't know themselves.

This is partially because values are symbolic constructs, concepts like fairness, freedom, dignity don't exist in nature, they're signifiers that structure our social reality and their meaning is never fully stable or determinate, because they exist in a web of relationships to other signifiers (Augusto & Campos, 2007). Values are also contradictory. For example, when weighing freedom against equality, maximising one constrains the other. The same is true of efficiency and meaning. These inherent contradictions aren't bugs to be resolved, they are constitutive tensions. We don't know what our values are because values exist in the symbolic order, which is itself inconsistent, there is no meta-level from which to resolve contradictions.

AI cannot navigate this. AI requires a determinate objective function, a consistent optimisation target and resolvable trade-offs, like Pareto frontiers or utility weights. Human values are indeterminate as meaning shifts with context, inconsistent as they contain genuine contradictions and unresolvable because the antagonism is constitutive, rather than accidental. So, for example, AI trained on human judgments about fairness will find no consistent function because fairness is contested and contradictory, different symbolic contexts activate different meanings and there is no true ideal of fairness underneath these conflicts.

We may also want to consider the relationship between the unconscious and self-opacity. We don't have transparent access to our own minds but have repressed desires and hidden motivations. For Lacan, the unconscious isn't a hidden chamber of the psyche, it's a structural effect of signification (Lacan, 1997). Meaning is never fully present to consciousness and language always means more/other than we intend. (Lacan, 1997) So, the claim *I don't know what I want* is the fundamental condition of subjectivity and asking someone what they want can't give you ground truth because their answer is itself a signifier that slides.

Finally, l'objet petit a (the object-cause of desire) further explains this mismatch. L'objet petit a, the unattainable object that causes desire to circulate, manifests when we don't want the object; we want to keep desiring (Lacan, 1997). Satisfaction of desire is actually traumatic. So, an AI that truly satisfied human desires would be horrifying. According to Žižek, we don't actually want our desires satisfied because satisfaction would mean the loss of desire which would mean the loss of subjectivity: we are constituted by our wanting (Žižek, 2008). Getting what we want leads to *jouissance*, painful enjoyment. Common stories, like the Midas touch or the monkey's paw illustrate that literal satisfaction of stated desire is catastrophic, we need the gap between what we say we want and what we get.

In terms of AI alignment, this reveals several deeply concerning potential scenarios that could emerge from this mismatch of ways of structuring reality between human and AI. There is the instance of literal interpretation, that is giving AI the command to make everyone happy, may lead to wireheading (Everitt et al., 2017), ignoring the fundamental complexities of the task. We may have technically got what we asked for, this scenario is horrifying because it bypasses desire's structure.

Also, our actual preferences being revealed to us through the missteps of AI is equally unfortunate. That is, AI optimises for what humans actually choose (e.g. clicks, purchases, engagement), which social media already does by giving us what we demonstrate we want (e.g. outrage, comparison, doomscrolling) rather than giving us what we may state as our preferences (e.g. fulfilment, purpose). The key insight here is that our revealed preference does not reflect what our choices indicate we really want. There is also the possibility of idealised volition, where AI gives us what we would want if we were wiser, but the fantasy of our ideal self is itself an objet a, and actualising it would destroy the fantasy. We need to remain imperfect and seeking. So, if we were able to reach perfect alignment this would still be dystopian because it would actually satisfy our desires and it is within the frustration of our desires that we have more hope of finding satisfaction.

This mismatch has major consequences for our prospects of reaching AI alignment. One reason this is key for alignment is that the Specification Problem runs deeper than we think. Revisiting Goodhart's Law, where the standard view is that when a measure becomes a target, it ceases to be a good measure (El-Mhamdi & Hoang, 2024), if we apply a psychoanalytic view, we can attest that human values are already internally contradictory, there's no true measure underneath. We don't want what we say we want; we want to remain unsatisfied (desire must remain unfulfilled). For example, people may claim to want efficiency but in fact we may actually need meaningful struggle to be fulfilled, actuated humans.

We may also consider the barrier that the impossibility of Coherent Extrapolated Volition (CEV) poses to alignment. CEV proposes AI should do what we would want if we knew more, thought faster, and were more the people we wished we were (Gabriel, 2020). However, the Lacanian subject is split, contradictory, structured by fantasy and there is no coherent "us" underneath. The concept of the people we wish we were is itself a fantasy object, not a real possibility. According to Žižek, ideology isn't false consciousness that hides reality; it structures reality itself (Žižek, 2008).

Moreover, reward hacking can be considered as an ontological, not a technical problem. For instance, we can reframe mesa-optimisation, where inner optimizers develop goals different from outer objectives (Skalse et al., 2022.), as resembling how human unconscious desire differs from conscious goals. AI's inner optimizer still lacks

the structure of subjectivity. Again, for Sycophantic AI, where models trained on human feedback learn to tell us what we want to hear, reveals that we don't actually want truth, we want validation instead. Yet we also want to want truth (desire at meta-level). AI cannot navigate this contradiction because it's not structured by desire.

We may also consider the impact of the anthropomorphism trap on AI alignment. Alignment research anthropomorphises AI by assuming AI can have goals or values in our sense while simultaneously de-subjectivising humans by treating human values as coherent utility functions. Both approaches rest on a fundamental mistake: humans aren't optimizers, and AIs aren't subjects. The gap isn't bridgeable by better specification.

In terms of AI safety, this has several implications.

Positive Implications	Negative Implications
We should abandon perfect alignment because trying to make AI "want what we want" is conceptually impossible	This means however, that some alignment approaches are doomed, such as CEV, value learning from human behaviour, etc.
We should focus on constraints, not goals by building systems with hard limitations rather than aligned objectives	The control problem remains because we still need to ensure AI doesn't cause harm, even if perfect alignment is impossible.
We should embrace human-in-the-loop and keep human subjects (with their contradictions) in the decision process	
We should recognise AI as tool, not an agent and not anthropomorphise AI systems into pseudo-subjects	

However, one might object that this equates to just saying alignment is hard, which we already knew. This rests on a misunderstanding, the argument is that it's impossible in principle because of an ontological mismatch. It's not that we need better specification methods; specification itself misunderstands human subjectivity.

Another objection may be that even if humans are contradictory, we can still satisfy surface preferences. However, surface preferences are often opposite of what we actually want. For example, giving people what they click on social media or giving them what actually fulfils them.

Furthermore, one may argue that this makes AI safety impossible. However, it redirects just the project, not eliminates it. Safety may be achieved through limitation rather than alignment, so acknowledging the problem is better than pursuing impossible solutions.

So, we may think of AI as a tool, not agent with values and safety as engineering problem (constraints) not philosophical one (alignment). As such, we ought to keep fundamental decisions in human hands because only subjects can navigate desire's contradictions and use AI for clearly specifiable tasks, not the pursuit of open-ended goals.

In conclusion, the psychoanalytic account of subjectivity reveals why the alignment problem is harder than anticipated, both technically and conceptually. Human wanting is structured by lack, mediated by the symbolic order, and irreducibly contradictory and AI systems cannot want in this sense. So perhaps the real danger isn't that AI won't want what we want, but that it will take what we ask for we want literally, thereby giving us a world optimised for our stated preferences, which would be a nightmare precisely because it wouldn't account for desire's necessary dissatisfaction. The solution isn't better alignment but the recognition that some problems (navigating contradictory desires, making choices in radical uncertainty) require subjects and cannot be delegated to optimisation processes gives us a starting point.

Augusto, P., & Campos, C. (n.d.). *The Lacanian Graph Hidden Behind Žižek's Project of Critique of Ideology*.

Christian, Brian. (2021). *The alignment problem : how can artificial intelligence learn human values?* Atlantic Books.

El-Mhamdi, E.-M., & Hoang, L.-N. (2024). *On Goodhart's law, with an application to value alignment*. <http://arxiv.org/abs/2410.09638>

Everitt, T., Krakovna, V., Orseau, L., Hutter, M., & Legg, S. (2017). *Reinforcement Learning with a Corrupted Reward Channel*. <http://arxiv.org/abs/1705.08417>

Gabriel, I. (2020). Artificial Intelligence, Values, and Alignment. *Minds and Machines* 2020 30:3, 30(3), 411–437. <https://doi.org/10.1007/s11023-020-09539-2>

Zizek, Slavoj. (2008). *The sublime object of ideology*. Verso.

Lacan, Jacques. (1997). *Ecrits : a selection*. W.W. Norton.

Russell, S. J. . (2020). *Human compatible : artificial intelligence and the problem of control*. Penguin Books.

Sbriglia, Russell., Žižek, Slavoj., Johnston, Adrian., Van Wert, Kathryn., Gorelick, Nathan., Rothenberg, M. Anne., Dolar, Mladen., Cole, Andrew., Radnik, Borna., McGowan, Todd., & Zupančič, Alenka. (2020). *Subject lessons : Hegel, Lacan, and the future of materialism*. Northwestern University Press.

Skalse, J., Howe, N. H. R., & Krueger, D. (n.d.). *Defining and Characterizing Reward Hacking*.

Žižek, Slavoj. (2013). *Less than nothing : Hegel and the shadow of dialectical materialism*. Verso.